

Embedded systems engineering

David Kendall

Northumbria University

- What's wrong with *super-loop*?
- What is a time-triggered scheduler?
- How to construct a schedule?
- What are the pros and cons?
- How can I implement a time-triggered scheduler in practice?
 - Based on Pont but for ARM target and with some restructuring of code.

Super loop architecture

```
/* My embedded system as a super loop */  
  
init();                /* Prepare to start          */  
  
while (true) {          /* repeat forever          */  
    f();                /* execute the control function */  
}
```

- Pros

- Simple – easy to understand
- Uses almost no resources

- Cons

- Difficult to ensure that $\tau()$ is called at precise instants of time.
- Many embedded systems require precise timing
 - Periodic tasks
 - One-shot tasks

Can we fix super loop?

```
/* My embedded system as a super loop */  
  
init ();                                /* Prepare to start */  
  
while (true) {                          /* repeat forever */  
    f ();                               /* execute the control function */  
    delay(n);                           /* delay for n microseconds */  
}
```

- This might work – repeat $f()$ every $m + n$ microseconds, where m is the execution time of $f()$
- But
 - to choose n we need to know m precisely
 - execution time of $f()$ must be the same each time round the loop
- Unrealistic assumptions

Fix number 2

```
/* My embedded system as a super loop */  
  
init();  
  
while (true) {  
    start = getCurrentTime();  
    f();  
    delay(start + p - getCurrentTime());  
}
```

- This is better
 - repeat $f()$ every p microseconds (more or less)
 - time for $f()$ can vary on each iteration but period remains constant
- But ...
 - Need to allow for time taken to get the time and configure the delay
 - Difficult to break controller into multiple functions that can execute at different rates

Towards a better solution

- Use timer-based interrupts to ensure that functions are called at precise instants of time.
- For example, ...

```
void f(void);                /* Control function prototype */

int main () {
    initTimer(TIMER0, f, 10); /* Set timer to interrupt at 10Hz */
    while (true) {
    }
}

void f(void) {                /* Control function – interrupt handler */
    ledToggle(USB_LINK_LED);

    /* clear interrupt */
    clearInterruptTimer(TIMER0);
}
```

Executing multiple tasks at different time intervals

- Embedded system may consist of multiple tasks that need to execute at different time intervals, e.g.
 - Read input from an ADC every millisecond
 - Read one or more switches every 200 milliseconds
 - Update LCD display every 3 milliseconds
- How to solve this problem?
 - Use multiple timers?
 - No - why not? ...coming next
 - Use a time-triggered scheduler?
 - Yes

Why not use multiple timers?

- May not have enough timers
 - e.g. 100 tasks into 4 timers does not go
- Code becomes hard to maintain
 - e.g. Change of oscillator frequency may involve modification to all tasks
 - e.g. adding another task not be possible if all timers are currently used
- Need to handle simultaneous interrupts
 - difficult to manage, hard to predict behaviour
 - system much simpler if there's only a single interrupt source

What is a time-triggered scheduler?

- extraordinarily simple operating system that allows tasks to be called on periodic and/or one-shot basis
- single timer ISR shared by many tasks, so
 - only one timer needs to be initialised
 - changes of timing source require only local code changes – usually one function at most
 - same scheduler can be used no matter how many tasks
- The time-triggered scheduler relies on a *static* schedule for its correct operation.

Static and Dynamic Scheduling

Static scheduling

In the *static scheduling* approach, all decisions about which task should run at any given time are made *offline*, i.e. *before* run-time. The job of the scheduler at run-time is very simple: it consults a scheduling table to see which task should run next and runs it. Typically, execution of a task is *non-preemptive*, i.e. it *runs to completion*.

Dynamic scheduling

In the *dynamic scheduling* approach, decisions about which task should run are made *online*, i.e. at run-time. The job of the scheduler is to determine which task should run next, according to some criteria, and then run it. Typically, execution of a task is *preemptive*. Examples of dynamic scheduling algorithms include fixed priority preemptive algorithms such as rate monotonic and deadline monotonic, and dynamic priority algorithms such as earliest deadline first and least laxity.

Periodic task model

- We assume that the job of the time-triggered scheduler is to run a set of periodic and one-shot tasks at pre-defined times.
- The periodic tasks are characterised by their
 - **period** (p)
 - **phase** also known as *offset*, (ϕ)
 - **worst-case execution time** (e)
 - **deadline** (d)
- We assume a system of N tasks comprises an indexed set of periodic tasks T

$$T = \{T_i : i \in 1..N\}$$

- Each periodic task can be regarded as generating *instances* of itself for execution, with instances numbered starting at 0.
- The **arrival time** of the j th instance of task i is

$$\alpha(T_{i,j}) = j * p_i + \phi_i$$

- **Harmonic periods** – the periods of a task set are *harmonic* iff every period in the task set is an integer multiple of all smaller periods in the set
- **Hyperperiod** – the *hyperperiod* of a task set is the greatest time that elapses until the pattern of task arrivals is repeated
 - The hyperperiod is equal to the least common multiple (LCM) of the periods of tasks in the set.
 - For a task set with harmonic periods, the hyperperiod is equal to the greatest of the periods of the tasks in the set.
- **Utilisation** – the *utilisation*, U , of a task set $T = \{T_i \mid i \in 1..N\}$ is given by

$$U = \sum_{i=1}^N \frac{e_i}{p_i}$$

Structured time-triggered scheduler

- A time-triggered scheduler can be implemented simply by using a *periodic* timer interrupt.
- The period of the timer interrupt defines a *frame* length
- Several tasks may be scheduled sequentially within a frame
- Let Z be the frame length and H be the hyperperiod. Then, the table that drives the scheduler has $F = \frac{H}{Z}$ entries
- Each entry lists the jobs to be executed in that frame
- The scheduler
 - is called by the timer interrupt
 - determines which frame should be scheduled
 - executes all jobs sequentially in the current frame
- The schedule repeats itself every hyperperiod.

Requirements for a schedule

- Hyperperiod H is least common multiple of periods of task set
- Frame size Z should be an integer divisor of H , at least as big as $\max\{e_i\}$ and no bigger than $\min\{p_i\}$. Usually $\gcd\{p_i\}$ is a good choice.
- Every job instance should be scheduled in exactly one frame
- No job instance should be scheduled before its release time
- The sum of the worst case execution times of the job instances scheduled in any frame should be no bigger than the frame size
- The deadline for any job instance should be no earlier than the start of the next frame following the one in which it's scheduled
- These requirements can be expressed formally, e.g. as an *integer linear program (ILP)*, and a schedule can be produced automatically by a solver for task sets of moderate size.

- Much of the (early) literature about embedded systems development introduces the ideas of time-triggered scheduling using a different vocabulary.
- The time-triggered scheduler is known as a **cyclic executive**
 - Manually constructed, off-line schedule of periodic tasks (procedure calls)
- Concurrent design, but sequential code (collection of procedures)
- Procedures are mapped onto a sequence of minor cycles (frames)
- Minor cycles constitute the complete schedule: the major cycle (hyperperiod)

Cyclic executive: properties and requirements

- No actual processes exist at run-time (only procedures)
- Minor cycles are sequences of procedure calls
- All periods must be a multiple of minor cycle time
- General rule:
 - minor cycle time is **gcd** of periods
 - major cycle time is **lcm** of periods
- Procedures share a common address space
 - Useful for inter-"process" communication
 - Only need one stack for user processes
 - No need for memory protection: concurrent access not possible
 - Deadlines are guaranteed by the offline schedule

Time-triggered scheduler: against – traditional view

- Difficult to incorporate:
 - processes with long periods
 - major cycle time determines maximum period
 - can (sometimes) be (partially) solved with secondary scheduling
 - processes with long computation times: must be split into several procedures
 - processes that are **sporadic** (not periodic but with well-defined minimum inter-arrival time)
- Difficult to construct and maintain the schedule
 - Fixed number of fixed sized procedures required
 - May cut across useful and well-established boundaries
 - Potentially very bad for software engineering (error prone)
- More flexible scheduling methods are difficult to support
- Determinism is an unnecessarily strong property; what is required is **predictability**

Time-triggered scheduler: case for the defence

- Generally accepted benefits:
 - The scheduler is simpler
 - The overheads are reduced
 - Testing is easier
 - Certification authorities tend to support this form of scheduling
- Most damaging of the problems:
 - Long-running computations either make the system unresponsive (minor cycle time too great) or must be split up artificially
- But Michael Pont [PON10, chp 13] claims:
 - In many systems, computations *are* extremely short
 - There are many sound techniques for decomposing long-running computations in practice
 - Increased micro-controller performance is reducing this problem
 - Where increased micro-controller performance is still not good enough, add more processors

Time-triggered scheduler: use in practice

```
int main() {
    schInit();           /* configure scheduler — 1 ms tick */

    schAddTask(f, 0, 5); /* prepare to run task f every 5 ms */
    schAddTask(g, 1, 10); /* prepare to run task g every 10 ms */
    schAddTask(h, 3, 15); /* prepare to run task h every 15 ms */

    schStart();          /* start ticking ... */

    while (true) {
        schDispatch();    /* if a task is ready to run, run it */
    }
}
```

Time-triggered scheduler: implementation

```
#ifndef __SCHEDULER_H
#define __SCHEDULER_H

#include <stdint.h>
#include <bsp.h>
#include <ttSchedConfig.h>

/* Task Control Block structure */
typedef struct schTCB {
    pVoidFunc_t task;
    uint32_t delay;
    uint32_t period;
    uint8_t invocations;
} schTCB_t;

void schInit(void);           // initialise the scheduler
void schStart(void);         // start ticking
void schUpdate(void);        // update after a tick — ISR
void schDispatch(void);      // run the next task
void schAddTask(
    pVoidFunc_t,              // the task to add
    uint32_t,                 // the delay in ms
    uint32_t);                // the period
void schRemoveTask(
    uint8_t);                 // remove a set from the task set
                             // identifier of the task to remove
void schSleep(void);         // go to sleep to save power

#endif
```

Acknowledgements

- Pont, M., [Patterns for Time-triggered embedded systems](#), TTE Systems, 2010
- Liu, J., Real-time systems, Prentice Hall, 2000