

# Embedded systems engineering

## Event-driven systems

David Kendall

Northumbria University

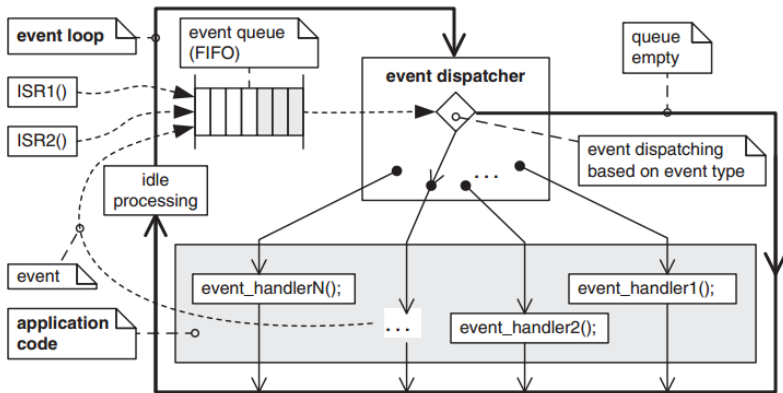
- How to manage communication of external events from ISRs to tasks
- How to manage communication between tasks
- How to ensure that the system remains responsive to all events for which it is responsible.
- One disciplined approach is to use an **event-driven framework**
- The approach presented here follows Chapter 6 of Samek, M., Practical UML Statecharts in C/C++ Event-Driven Programming for Embedded Systems, 2nd Edition, Newnes, 2008

- Samek introduces a *framework* for event-driven embedded systems that is intended to support a development approach based on *hierarchical state machines* (HSMs)
- His approach is useful *whether or not* you intend to develop with HSMs
- What you get is a *disciplined* approach to event-driven system development that produces code that is:
  - free from concurrency hazards
  - responsive
- The rest of the lecture often uses state machine examples but the approach is more generally applicable

# Inversion of control

- Traditional sequential programs have a control model in which the program waits (blocks) for an event (input) whenever it needs it
  - the programmer manages the control code, but
  - the program is *unresponsive* to other events while waiting
- Traditional approach to *multi-tasking* was developed so that programmers could maintain the blocking approach to I/O with which they were familiar, while allowing the program to remain responsive
  - while one or more tasks are waiting for I/O, run some other task that is ready
  - Disadvantage: task overheads: one stack per task, context switch etc.
- Event-driven systems *invert the control*
  - **Hollywood principle** - “don’t call us, we’ll call you”
  - the event-driven system manages the control code
  - the programmer writes the event handlers

# Traditional event-driven system



(Samek, 2008, p.264)

# Traditional event-driven system (TEDS)

- The traditional event-driven system is clearly divided into
  - the event-driven infrastructure
    - event loop
    - event dispatcher
    - event queue
  - the application
    - event handlers
- Events can be generated by ISRs or by event handlers
- All events go into the event queue
- Event dispatcher is in control of the system
  - remove event from event queue
  - dispatch to appropriate event handler based on the *event type*
- Event handler - simple, non-blocking, sequential code that runs to completion and returns to the dispatcher as soon as possible
- Widely-used approach in GUI frameworks, e.g. MFC, X-Windows, Swing, Qt etc.

# Characteristics of TEDS

- Pro

- Flexible patterns of control
- More efficient use of the CPU than in a traditional sequential system
- Typically consumes less stack space than a traditional multi-tasking system

- Con

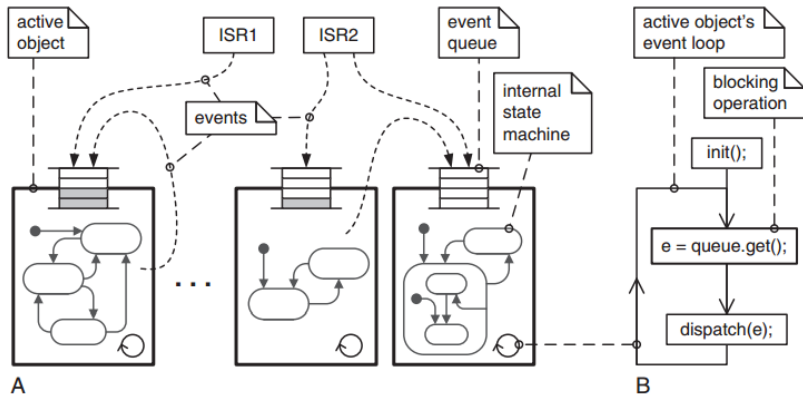
- Responsiveness could be better - single event queue makes it difficult to prioritise work
- Encapsulation difficult - event handlers must remember their state by using global variables

# Active object computing model

- Active object computing model addresses the problems of TEDS
- Based on the “actor” model developed by Carl Hewitt and colleagues in the 70s and 80s
  - autonomous software objects communicating by message-passing
- UML specification includes the concept of active object
- Essential idea
  - Use multiple event-driven systems in a multi-tasking environment
  - Or as Samek puts it
    - Active object = thread of control + event queue + state machine



# Active objects system



(Samek, 2008, p.267)

# Characteristics of the active object computing model

- Asynchronous communication
  - Active objects receive their events exclusively via their event queues
  - Event producers post events to event queues but don't wait for response
  - No distinction between events generated by ISRs and events generated by active objects
- Run-to-completion
  - RTC semantics guaranteed by structure of each AO's event loop - `dispatch()` must complete and return before next event can be extracted from event queue
  - Notice that here RTC is *not incompatible* with a preemptive RTOS
    - RTC step can be preempted by another AO's thread of control (task) without any concurrency hazards, assuming that *AO's do not share resources*

# Characteristics of the active object computing model

- Encapsulation

- All data and other resources are *local* to the active object, i.e. active objects have *no shared resources*
- Communication with the outside world and with other active objects is restricted to asynchronous event exchange
- Event exchange and the event queue are managed by the framework and are guaranteed to be free from concurrency hazards
- Encapsulation does not require an object-oriented language - just a disciplined approach to *information hiding*

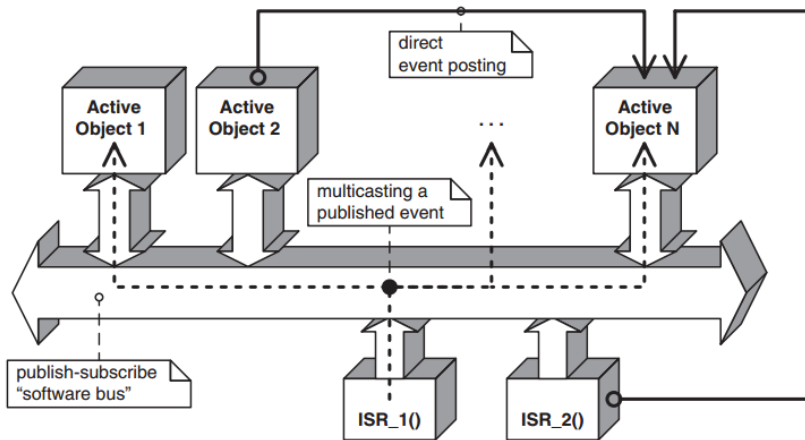
- Complements traditional multi-tasking RTOS

- Most common implementation of active objects is to map them to tasks of a traditional preemptive RTOS, e.g. uCOS-II
- But the active object model can also be implemented using a very simple cooperative (non-preemptive) scheduler

# Event delivery mechanisms

- A key component of an event-driven framework is the *event delivery mechanism*
  - Responsible for the efficient transfer of events from producers to consumers
- Two common types of delivery:
  - *Direct event posting*: producer of the event puts it directly into the event queue of the consumer
  - *Publish-subscribe*: producer “publishes” events to the framework, the framework delivers the event to the event queues of all active objects that “subscribed” to it

# Event delivery mechanisms



(Samek, 2008, p.280)

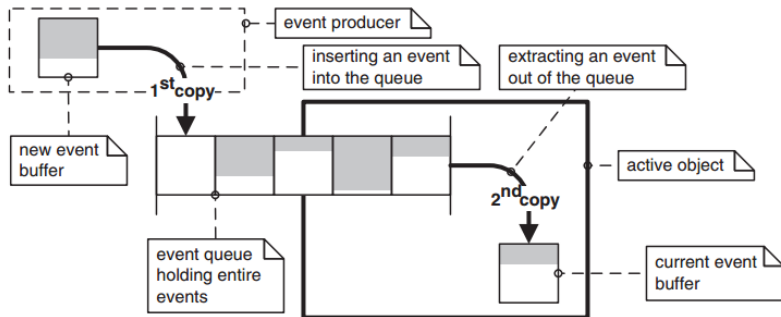
# Event delivery mechanisms

- Direct event posting
  - Simple to implement
  - Push-style communication: active objects receive events whether they want them or not
  - Tight coupling between producers and consumers
- Publish-subscribe
  - Producers and consumers are loosely coupled
  - A “mediator” is required to accept published events and deliver them to subscribers
  - The mediator must be capable of storing information about subscribers, allowing dynamic subscription and cancellation
  - Efficient, thread-safe multicasting of events is required

# Event memory management

- Events are produced and consumed *dynamically* (and frequently)
- Dynamic production and consumption of events is the key function of an event-driven system
- Efficient management of the memory used by events is crucial, since this memory must be reused as new events are produced and old events are consumed
- A challenge for the event-driven framework is to ensure that event memory is not reused until all active objects have finished their RTC processing of the event

# Event memory management - copy



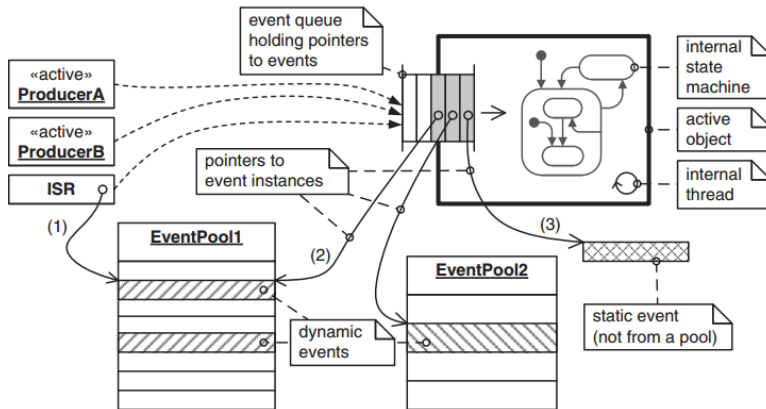
(Samek, 2008, p.283)



# Event memory management

- This approach is *safe* - avoids corruption of event memory before the event has been processed
- But it's *expensive* - copying an entire event from the producer's event buffer into the consumer's event queue and then copying from the event queue into the consumer's event buffer introduces a significant overhead
- For this reason, it is common to store just a *pointer* to the event memory in the event queue
- The main problem with this approach is that the producer of the event needs to know when all consumers have finished with it before it reuses the event memory
- An event-driven framework can maintain control of the entire event lifecycle and so can use an efficient pointer-passing approach to event exchange and arrange for the safe *garbage collection* of event memory once all active objects have processed the event

# Event memory management - zero-copy

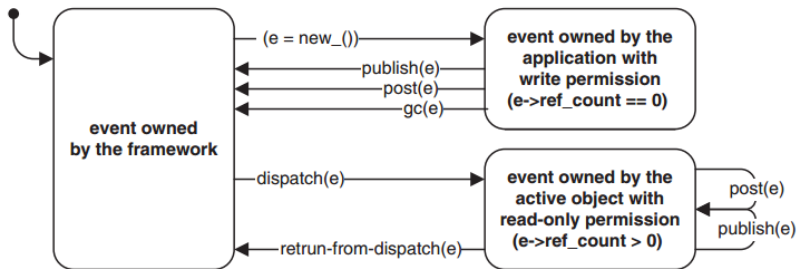


(Samek, 2008, p.285)

# Event memory management

- uCOS-II provides exactly the right kind of features to make the management of event memory straightforward
- It offers *message queues* with thread-safe *pend* and *post* operations for pointer-sized data
- It offers *memory pools* for the allocation and deallocation of fixed size blocks of memory - avoids memory fragmentation (problem with malloc and free) ; offers deterministic performance
- Event-driven framework can implement a simple *reference counting* garbage collection algorithm to handle the multicasting of events in a publish-subscribe system
- However, the application code must follow a disciplined approach to *event ownership*

# Event ownership



(Samek, 2008, p.288)

# Acknowledgements

- Samek, M., Practical UML Statecharts in C/C++ Event-Driven Programming for Embedded Systems, 2nd Edition, Newnes, 2008