

Embedded Systems Engineering

System Reliability - 2

- Application software aspects
 - Basic design issues
 - Run-time problems
- Real-world interfacing
 - Identifying & detecting faults on inputs
 - Avoiding faults on outputs
- Operating systems aspects
 - Protection techniques
 - Tasking models & scheduling
 - Distributed applications

Michael Brockway

Application Software Aspects

- Basic program design issues
 - Use rigorous design techniques
 - Use well-ordered program structures
 - Develop & use good programming standards
 - Write *readable* code
 - Use an appropriate programming language. For critical applications -
 - ◆ Spark Ada, Ada 95, MISRA C
 - Use code quality checking tools
 - ◆ Lint for C
 - Do not use unconditional transfers of program control
 - Do not use recursion
 - Avoid pointers
 - Limit use of inheritance, polymorphism
 - For highly critical systems, do not use dynamic memory allocation
 - Use static & dynamic code analysis techniques

Application Software Aspects

- Dealing with run-time problems
 - Exception handling
 - Backward error recovery
 - N-version programming
- Exception Handling
 - Normal operation ceases: control transfers to an *exception handler*
 - Subsequent events are application-specific
 - ◆ determined by functional needs, response times, system criticality
 - It may be possible to put the system into a pre-determined acceptable state and then continue processing
 - ◆ *Forward error recovery*

Exception Handling

- Many languages do not have this. Possibly the neatest implementation is in Java: eg
 - The `OutputStream` constructors and the functions `writeObject`, `close` may **throw** an `IOException`:
 - Object-oriented: exceptions are encapsulated in objects of subclasses of class `Exception`

```
ObjectOutputStream str;  
try {  
    str = new ObjectOutputStream(new FileOutputStream(  
        chooser.getSelectedFile().getName()));          //open file  
    Enumeration e = strokes.elements(); //Iterate thru strokes,  
    while (e.hasMoreElements()) {                      //writing to file.  
        str.writeObject(e.nextElement());  
    }  
    str.close();                                       //close file  
} catch (IOException x) {  
    System.err.println("Could not open file for output");  
    x.printStackTrace();  
}
```

Exception Handling

- C++ has try/throw/catch too. But generally the programmer decides when an exception will be thrown

```
int getNextNum() {
    int n;
    ifstream.input("nums.txt");
    if (!input) throw "Input Failed";
    input >> n;
    input.close();
    n++;
    ofstream.output("nums.txt");
    if (!output) throw "Output Failed";
    output << n;
    output.close();
    return n;
}

void main() {
    try {
        cout << getNextNum() << endl;
    }
    catch (char * error) {
        cout << "Error: " << error << endl;
    }
    //processing resumes from here
}
```

Exception Handling

- C has the **assert** macro -

```
double a, b, c, s, discr;  
...  
discr = b*b - 4*a*c;  
assert(discr >= 0);  
s = sqrt(discr);  
printf("Solutions are %f, %f\n", (-b+s)/2/a, (-b-s)/2/a);  
...
```

- If the assertion fails, the program is aborted with a diagnostic message on the console.
 - Uses the **abort** facility

Exception Handling

- C also has the **signal** facility -

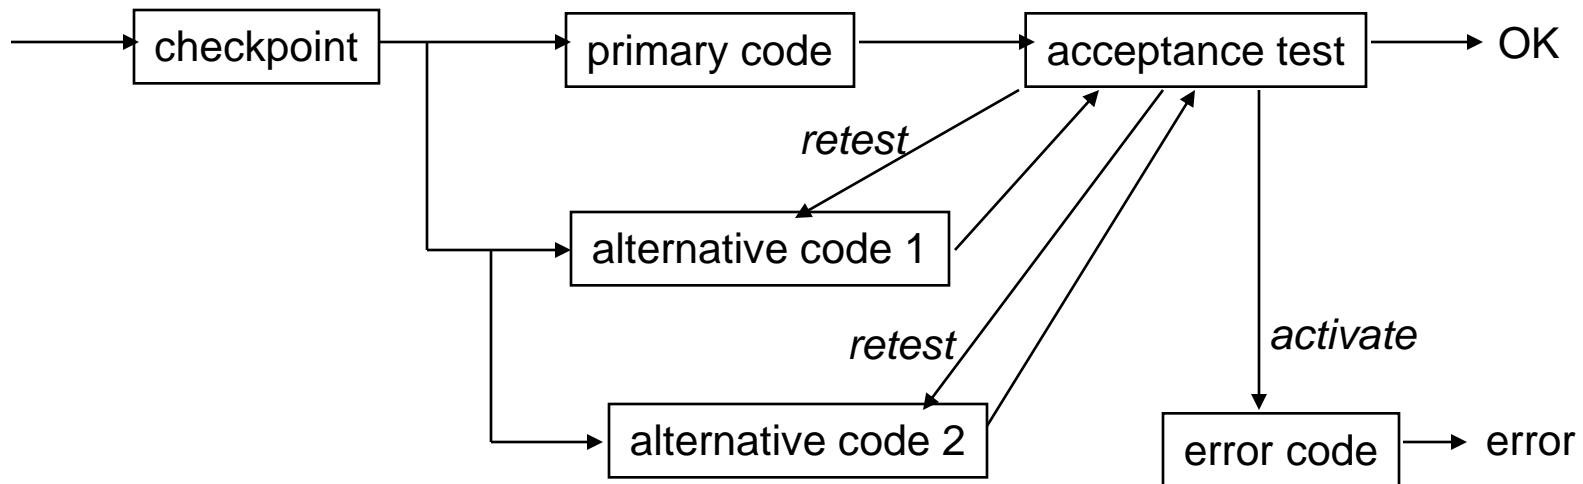
```
void myHandler(int theSignal) { ... }

int sig;
void (* oldHandler)();
/* set new handler, saving old handler */
oldHandler = signal(sig, &myHandler);
if (oldHandler == SIG_ERR)
    printf("Could not establish new handler\n");
```

- A signal can be raised by the computer's error detection mechanisms, or by a program with `raise(int sig);`
- Pre-defined values of `sig` which might be raised by the computer include
 - SIGABRT – abnormal termination (abort facility)
 - SIGFPE – floating point error or divide-by-0
 - SIGILL – invalid instruction
 - SIGSEGV – invalid memory access
 - SIGTERM – termination signal from a user or another program

Backward Error Recovery

- Aka *rollback*
- Maintains continuous operation when a failure occurs
- A common method is that of *recovery blocks*:



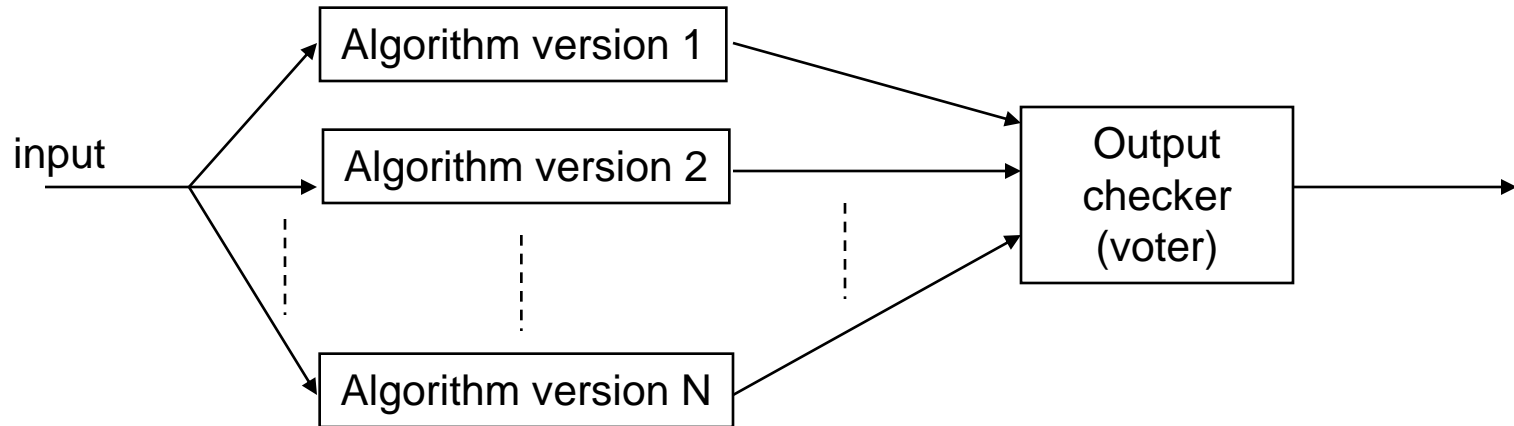
Backward Error Recovery with Recovery Blocks

- System state saved at checkpoint, then primary algorithm executed.
- If acceptance test is passed, processing continues
- Otherwise,
 - processing is *rolled back* to the state at the checkpoint
 - processing resumes with alternative code 1
 - acceptance test applied again
- On another failure,
 - roll back again
 - alternative code 2
- etc
- If all alternatives tried and acceptance test still fails,
 - execute error code

Backward Error Recovery with Recovery Blocks

- This approach is more suited to mission-critical than safety-critical systems because
 - A program failure may cause operation to stop completely until some form of external recovery is put into action;
 - Actual execution times may vary from run to run
 - ◆ performance problems
 - ◆ hard deadlines could be missed
 - ◆ processing time is indeterminate
- Designing the Acceptance tester – 3 approaches -
 - test results against pre-defined values
 - test results against predicted values
 - ◆ eg from advance knowledge of maximum rate of change
 - determine all the input values which could have produced this output: compare with checkpoint value(s)
 - ◆ can be time consuming


N-version Programming



- Checker
 - compares all results
 - If all agree, outputs the result; otherwise
 - selects the result by majority & outputs it.
 - ◆ Normally N-1 agreements at any one time
- Roots in electronic analogue control systems but can be used in single-processor applications

N-version Programming

- No need to devise acceptance tests
 - Actual results do not matter provided we have majority agreement
 - But there is time overhead



time →

- Fails when common-mode errors occur
 - all versions give wrong result
 - Minimise by using as much diversity as possible between the algorithms – each developed by a different designer

Real-world Interfacing

- Lutz, working on Voyager & Galileo spacecraft found poor understanding of interfacing requirements accounted for 44% of all logged safety-related errors -
 - Out-of-range input values
 - Non-arrival of expected inputs
 - Unexpected input arrival
 - Inconsistent code behaviour in response to input signals
 - Invalid input data time frames
 - Out-of-range arrival rates
 - Lost events
 - Excessive output signal rates
 - Not all output data used
 - Effect of input signal arrival during non-operational mode
 - ◆ start-up, off-line, shut-down

Real-world Interfacing

- Input problems –
 - switch inputs stuck in one state
 - uncommanded change of switch state
 - sensor signals going hard over to max/min
 - analogue signals locking up
 - invalid signals due to noise
 - bias on digitised signals due to stuck bits
 - sensor drift with time
- Input fault detection methods
 - Limit testing
 - ◆ Compare actual values with known practical limits.
 - ◆ Detects mainly the "hard over" type of fault
 - Rate testing
 - ◆ Compare rates of change of actual values with known practical maxima.
 - ◆ Extracting the rate of change takes time
 - ◆ differentiation magnifies effects of noise
 - ◆ subject to "false alarm" errors

Real-world Interfacing

- Input fault detection methods (ctd)
 - Predicted values
 - ◆ Noise can be a problem
 - ◆ A mathematical model of the system can be used to provide the predicted values
 - Redundant inputs
 - ◆ multiple versions of inputs cross-checked
 - ◆ In critical systems, dual redundancy common; sometimes triple, quad.
 - ◆ Disagreement does not identify the fault
 - ◆ but suitable in a system which can revert to a manual mode
 - ◆ Where continuous operation needed, use majority voting
 - Time-related values
 - ◆ Ignore signals outside specific times
 - ◆ When synchronizing operations, say, over a LAN, time-stamping messages is critical.

Real-world Interfacing

- Input fault detection methods (ctd)
 - Inferred values
 - ◆ An error on an input can be inferred from values of other inputs: eg an air-conditioning unit with sensors for room temperature, unit inlet and unit outlet temperature:
 - ◆ room = 20°, inlet = 0°, outlet = 10°
 - ◆ the inlet temperature sensor is the culprit
 - ◆ Can reduce redundancy of inputs needed in safety-critical systems
 - Estimated values
 - ◆ similar to model-based prediction; but a model may not be known in advance
 - ◆ Instead, the model is built "on the fly"

Real-world Interfacing

- Avoiding Faults on Outputs
 - Do not use individual bits within a word to operate separate on/off controls. Use a specific, unique word for each control
 - If each bit of a single word does operate an individual control, then
 - ◆ Store the word in memory
 - ◆ You can read it and see what state the outputs are *supposed* to be in
 - ◆ When setting outputs, write to the control word, then output this
 - ◆ Use AND, OR bit masking to ensure you do not affect other controls
 - Use state, sequence information to limit the number of operations that may be invoked
 - Use multiple signalling where operating a control incorrectly could be dire
 - ◆ should be tied into hardware so that proper status signals are obtained
 - Apply rate-limiting to analogue signals

Operating Systems Aspects

- Problems include
 - Application software interfering with the OS
 - Tasks/applications interfering with other tasks/applications
 - Unexpected functional behaviour of tasks/applications
 - Unexpected timing behaviour of tasks/applications
 - Design weaknesses in OS
- Dealing with interference – Protection techniques
 - Memory protection, to prevent writing to the data or code area of other tasks
 - Intertask communication, signalling to implement synchronisation or mutual exclusion
 - ◆ binary or counting semaphores
 - ◆ event flags
 - ◆ messages

Tasking Models and Scheduling

- Additional requirements for safety-critical systems
 - functionality is fully predictable
 - timing is fully predictable
 - guaranteed detection of disk, OS failures
 - require a small (& predictable) amount of memory
- At the *catastrophic* severity level, the following rules are a guide:
 - static task schedule (no dynamic creation/deletion of tasks)
 - All code can be statically analysed
 - Tasks run to completion, without pre-emption
 - Watchdog techniques ensure run-time bounds are not transgressed

Tasking Models and Scheduling

- At the *critical* severity level, the following rules are a guide:
 - Both periodic and aperiodic tasks are allowed
 - Fixed-priority tasks are supported; dynamic adjustment of priorities is forbidden
 - Non-pre-emptive, co-operative and pre-emptive scheduling are allowed
 - Schedule is static – no dynamic task creation/deletion
 - Periodic tasks are structured as infinite loops
 - Worst-case execution times are deterministic
 - Memory for kernel components – stacks, TCBs, ... – is allocated statically
 - All code can be statically analysed
 - No priority inversion
 - Use scheduling algorithms which lend themselves to schedulability analysis
 - Watchdogs for run-time bounds

Distributed Applications

- The system comprises a number of computing nodes linked by a network.
- Each node runs autonomously, but co-operates with other nodes by passing messages through the network
- A good design features
 - Composability
 - ◆ A large system is built by integrating a set of well-specified and tested subsystems, whose essential properties (timeliness, stability, ...) are preserved by the system integration.
 - Scalability
 - ◆ The system can grow by the addition of nodes, within the capacity of the communication system, or by the replacement of a node by a gateway to another communication network supporting additional node.
 - ◆ In this architecture, complexity is kept with reasonable bounds by encapsulation.

Distributed Applications

- a good solution when the application -
 - controls physical devices that are physically dispersed
 - ◆ plant control systems
 - controls physical devices that are dynamically interchangeable or configurable
 - ◆ train control systems
 - needs to be highly reliable and fault-tolerant
 - ◆ This can be provided by providing redundancy -- replicating components.
 - ◆ In a good design, an error is localised within a relatively simple local subsystem

Distributed Applications

- Possible problems for the developer -
 - The system depends on an effective communication network. This might mean there is a single point of failure
 - Performance overhead
 - ◆ A local procedure call might take 10 to 100 μ sec (microseconds)
 - ◆ A remote procedure call might take 10 or 20 msec (milliseconds) -- 200 to 1000 times as long.
 - In a *hard* real-time system, a result must be delivered within a specified time frame; otherwise the system has failed (ABS braking system, nuclear reactor control system). How does the developer *predict* real-time performance without being unduly pessimistic, when communication through a network is involved?

Distributed Applications

- Possible problems for the developer –
 - A failure could be
 - ◆ in the network
 - ◆ in a node CPU
 - ◆ in a task at a node
 - Data *addressing* errors (wrong senders/receivers)
 - Data errors (corruption)
 - Message timing errors
 - Message sequence errors
 - Consistency of data across the system
 - Synchronisation of operations across the system